

Homework 2 – SF2568 Parallel Computations for Large-Scale Problems

Cuong Duc Dao, Donggyun Park
cuongdd@kth.se, donggyun@kth.se

March 02, 2021

Question 1: The broadcast operation is a one-to-all collective communication operation where one of the processes sends the same message to all other processes.

- (a) Assume our simple communication model for point-to-point communication. A straightforward implementation would require $P - 1$ communication steps. Design an algorithm for the broadcast operation using only point-to-point communications which requires $O(\log P)$ communication steps.
- (b) Do a time performance analysis for your algorithm
- (c) How can the scatter operation be implemented using $O(\log P)$ communication steps?

Solution:

- (a) Using recursive doubling we divide the data distribution according to the Figure 1. It only requires $O(\log_2 P)$. If we set the range of an arbitrary number of processors P for $2^T \leq P < 2^{T+1}$, the pseudo-code for this algorithm is as below 1

Algorithm 1 Recursive doubling algorithm for broadcasting

```
1: if  $p \leq 2^T$  then
2:   send(data, bitflip(p, T))
3: end if
4: if  $p < P - 2^T$  then
5:   receive(data, bitflip(p, T))
6: end if
7: if  $p < 2^T$  then
8:   for  $t = 0$  to  $T-1$  do
9:     loc  $\leftarrow$  bitflip(p, t)
10:    if  $p < 2^t$  then
11:      send(data, loc)
12:    else
13:      if  $p < 2^{t+1}$  then
14:        receive(data, loc)
15:      end if
16:    end if
17:  end for
18: end if
19: if  $p < P - 2^T$  then
20:   send(data, bitflip(p, T))
21: end if
22: if  $p \geq 2^T$  then
23:   receive(data, bitflip(p, T))
24: end if
```




Figure 1: Illustration of Recursive Doubling for broadcast operation between 8 processes

- (b)
- Local computation time: $t_{comp,1} = 1t_a$ (where t_a is the time for one basic operation)
 - Recursive doubling step: $t = T(t_{startup} + w_p t_{data} + 3t_a) = \log P(t_{startup} + w_p t_{data} + 3t_a)$ where $3t_a$ is from the two if statement evaluations and the computation of the destination. We have $t_{startup} + w_p t_{data}$ from the send or receive that is triggered at most once per loop.
 - Total computation time: $T_p = t_{comp,1} + t = t_a + \log P(t_{startup} + w_p t_{data} + 3t_a)$
- (c) Assuming that we have a message of length M which can be divided into P processors, so each processor would own a chunk of data of length M/P . The scatter operation sends each chunk to its corresponding processor. This operation can be done in $O(\log P)$ time steps using the **recursive doubling** strategy mentioned above. Processor with lower rank sends the upper half of the data to the middle node by performing the reverse bit-flip operation, i.e., from highest bit to lowest bit. A graphical illustration of this algorithm is given in Figure 2




Figure 2: Illustration of scatter operation with recursive doubling with $P = 8$

Question 2: Consider a matrix A distributed on a $P \times P$ process mesh. An algorithm has been given in the lecture for evaluating the matrix-vector product $y = Ax$. While x is column distributed, y is row distributed. In order to carry out a further multiplication Ay , the vector y must be transposed (the row-distributed y must be redistributed to become column-distributed).

- (a) Design an algorithm for this transposition. (1)
- (b) Make a performance analysis of your transposition algorithm. (1)
- (c) An extra credit will be given for a good (!) solution in the case that A is distributed in a $P \times Q$ process mesh with $P \neq Q$.

Solution:

(a) Let us assume that a vector y has N elements. It needs to be divided into the number of processors to be used in $P \times P$ process mesh. Then N/P elements are assigned for each processor. 'MPI_Alltoall' command distributes the data stored in a processor to every processor including itself. So when the data list and the processor aligned in the corresponding order, it automatically does the matrix transposition. Conceptually, it is the same as the scattering strategy Question 1, where each processor scatter its data to all the other processors.




Figure 3: This example shows how a vector has been divided into P chunks ($P = 5$ in this case) and each chunk has P elements and transposed by MPI_Alltoall command

After this transposition is done y is column-distributed, if necessary, we need to concatenate them back together, then we obtain the column distributed vector y .

(b) Because each processor can perform the scatter operation at the same time with other processors, the time complexity of this transposition is the same as the complexity of the scatter operation above: $T_p = t_{comp,1} + t = t_a + \log P(t_{startup} + wt_{data} + 3t_a)$

(c) Let's take an example with 20 elements.

In this case, $P = 4$ and $Q = 3$. Let's think about it with the given example below. We have this 4 by 3 matrix.

```
1  2  3
1  2  3
1  2  3
1  2  3
```

and we add 'NULL' (empty) at the end of each row. Then it becomes $P \times P$ (4x4 in this case). The data distribution would look like below.

```
1  2  3  NULL
1  2  3  NULL
1  2  3  NULL
1  2  3  NULL
```

Now we need to transpose as what we do in part (a). After finishing the transposition, it should be in the form below

1	1	1	1
2	2	2	2
3	3	3	3
NULL	NULL	NULL	NULL

Finally we have obtained the transposed version of the original.

Question 3: Approximate solution for 1D-poisson partial differential equation with Jacobi iterations.

Solution:

The data is linearly distributed over processes. The challenge is that each process would need data from some indexes (in the data vector) from its neighbour processes. This problem is tackled with red-black communication. The details implementation of the red/black algorithm can be found in the code listing (line 96 - 110).

Figure 4 and 5 present the results of our code for approximating the solution for the Poisson partial differential equation with Jacobi iterations. We can see that there is no notable difference between the approximated and the true function in the plots. Therefore, we can assume that the numerical approximation has reached convergence before 10^6 iterations.




Figure 4: $u(x) = x(x-1)$; $r(x) = x-1$; $f(x) = 2 + x(x-1)^2$




Figure 5: $u(x) = x(x - \frac{1}{2})(x - 1)$; $r(x) = x - 1$; $f(x) = 6x - 3 + (x - 1)x(x - \frac{1}{2})(x - 1)$

Implementation of Jacobi iteration method for approximating the solution for Poisson partial differential equation

```

1  /* Reaction-diffusion equation in 1D
2  * Solution by Jacobi iteration
3  * simple MPI implementation
4  *
5  * C Michael Hanke 2006-12-12
6  */
7  #include <stdlib.h>
8  #include <stdio.h>
9
10 #define MIN(a,b) ((a) < (b) ? (a) : (b))
11
12 /* Use MPI */
13 #include "mpi.h"
14
15 /* define problem to be solved */
16 #define N 1000 /* number of inner grid points */
17 #define SMX 1000000 /* number of iterations */
18
19 /* implement coefficient functions */
20 extern double r(const double x);
21 extern double f(const double x);
22
23 double r(const double x) {
24     // Change here before submission: -x
25     // return x - 1;
26     // return -x;
27     return x - 1;
28 }
29
30 double f(const double x) {
31     // return -2 + x * x * (x - 1);
32     // return 2 + (x - 1) * x * (x - 1);
33     return 6*x - 3 + (x-1)*x*(x-0.5)*(x-1);
34 }
35
36 /* We assume linear data distribution. The formulae according to the lecture
37 are:
38     L = N/P;
39     R = N%P;
40     I = (N+P-p-1)/P;    (number of local elements)

```

```

41     n = p*L+MIN(p,R)+i; (global index for given (p,i)
42     Attention: We use a small trick for introducing the boundary conditions:
43     - The first ghost point on p = 0 holds u(0);
44     - the last ghost point on p = P-1 holds u(1).
45     Hence, all local vectors hold I elements while u has I+2 elements.
46 */
47
48 int main(int argc, char *argv[])
49 {
50     /* local variable */
51     int P, p, tag, L, R, I, n, color, step, i;
52     double h, *u, *unew, *ff, *rr, x_n;
53     MPI_Status status;
54
55     /* Initialize MPI */
56     MPI_Init(&argc, &argv);
57     MPI_Comm_size(MPI_COMM_WORLD, &P);
58     MPI_Comm_rank(MPI_COMM_WORLD, &p);
59     if (N < P) {
60         fprintf(stdout, "Too few discretization points...\n");
61         exit(1);
62     }
63     tag = 100;
64     color = p % 2; // 0: red, 1: black
65     printf("Node %d, color: %d\n", p, color);
66     /* Compute local indices for data distribution */
67     L = N / P;
68     R = N % P;
69     I = (N + P - p - 1) / P; // Number of local elements
70     h = 1.0 / (N + 1);
71
72     printf("Processor %d, L = %d; R = %d; I = %d; h = %f\n", p, L, R, I, h);
73     /* arrays */
74     unew = (double *) malloc(I*sizeof(double));
75     /* Note: The following allocation includes additionally:
76     - boundary conditins are set to zero
77     - the initial guess is set to zero */
78     u = (double *) calloc(I+2, sizeof(double));
79
80     ff = (double *) malloc(I * sizeof(double));
81     rr = (double *) malloc(I * sizeof(double));
82
83     for (i = 0; i < I; i++) {
84         n = p * L + MIN(p, R) + i; // Global index for given (p, i)
85         x_n = n * h;
86         ff[i] = f(x_n);
87         rr[i] = r(x_n);
88     }
89
90     /* Jacobi iteration */
91     for (step = 0; step < SMX; step++) {
92         if (step % 100000 == 0) {
93             printf("Jacobi, processor %d, iter %d\n", p, step);
94         }
95         /* RB communication of overlap */
96         if (color == 0) { // red
97             MPI_Send(&u[I], 1, MPI_DOUBLE, p+1, tag, MPI_COMM_WORLD);
98             MPI_Recv(&u[I+1], 1, MPI_DOUBLE, p+1, tag, MPI_COMM_WORLD, &status);
99             if (p != 0) {
100                 MPI_Send(&u[1], 1, MPI_DOUBLE, p-1, tag, MPI_COMM_WORLD);
101                 MPI_Recv(&u[0], 1, MPI_DOUBLE, p-1, tag, MPI_COMM_WORLD, &status);
102             }
103         } else { // black
104             MPI_Recv(&u[0], 1, MPI_DOUBLE, p-1, tag, MPI_COMM_WORLD, &status);
105             MPI_Send(&u[1], 1, MPI_DOUBLE, p-1, tag, MPI_COMM_WORLD);
106             if (p != P-1) {
107                 MPI_Recv(&u[I+1], 1, MPI_DOUBLE, p+1, tag, MPI_COMM_WORLD, &status);
108                 MPI_Send(&u[I], 1, MPI_DOUBLE, p+1, tag, MPI_COMM_WORLD);
109             }
110         }
111         /* local iteration step */
112         for (i = 0; i < I; i++) {
113             unew[i] = (u[i]+u[i+2]-h*h*ff[i])/(2.0-h*h*rr[i]);

```

```

114         // unew[i] = (u[i-1] + u[i+1] - h*h*ff[i]) / (2.0 - h*h*rr[i]);
115     }
116     for (i = 0; i < I; i++) {
117         u[i+1] = unew[i];
118     }
119 }
120
121 /* output for graphical representation */
122 /* Instead of using gather (which may lead to excessive memory requirements
123 on the master process) each process will write its own data portion. This
124 introduces a sequentialization: the hard disk can only write (efficiently)
125 sequentially. Therefore, we use the following strategy: */
126 FILE *f;
127 double signal;
128 // 1. The master process writes its portion. (file creation)
129 if (p == 0) {
130     f = fopen("possession.txt", "w");
131     for (i = 0; i < I; i++) {
132         fprintf(f, "%f ", unew[i]);
133     }
134     fclose(f);
135     // 2. The master sends a signal to process 1 to start writing.
136     signal = 1.0;
137     MPI_Send(&signal, 1, MPI_DOUBLE, 1, tag, MPI_COMM_WORLD);
138     printf("Master processor wrote and sent signal\n");
139 } else {
140     // 3. Process p waits for the signal from process p-1 to arrive.
141     MPI_Recv(&signal, 1, MPI_DOUBLE, p-1, tag, MPI_COMM_WORLD, &status);
142     printf("Received signal %f from processor %d\n", signal, p-1);
143     if (signal == 1.) {
144         // 4. Process p writes its portion to disk. (append to file)
145         f = fopen("possession.txt", "a");
146         for (i = 0; i < I; i++) {
147             fprintf(f, "%f ", unew[i]);
148         }
149         fclose(f);
150         printf("Processor %d wrote its portion\n", p);
151         // 5. process p sends the signal to process p+1 (if it exists).
152         if (p != P - 1) {
153             MPI_Send(&signal, 1, MPI_DOUBLE, p+1, tag, MPI_COMM_WORLD);
154             printf("Processor %d sent signal to processor %d\n", p, p+1);
155         }
156     }
157 }
158
159 /* That's it */
160 MPI_Finalize();
161 exit(0);
162 }

```
